

Bayesian Learning

Lab 1

Emil K Svensson and Rasmus Holm

2017-04-11

Question 1

In this assignment we are given data coming from a Bernoulli distribution with unknown parameter θ and we use the conjugate prior

$$\theta \sim \text{Beta}(\alpha_0, \beta_0)$$

where $\alpha_0 = \beta_0 = 2$. The sample we get consists of 14 successes out of 20 trials which gives us the posterior distribution

$$\theta|y \sim \text{Beta}(\alpha_0 + 14, \beta_0 + 6).$$

```
n <- 20
s <- 14
f <- n - s

alpha_prior <- 2
beta_prior <- 2
```

a

Here we are sampling from the posterior distribution to verify that as the number of samples increases, the mean and standard deviation converges to the true values. In this particular case the true mean is $\alpha/(\alpha + \beta) = 2/3$ and standard deviation is $\sqrt{(\alpha\beta)/((\alpha + \beta)^2(\alpha + \beta + 1))} = \sqrt{2/225}$.

```
drawSamples <- function(nDraws, alpha, beta) {
  rbeta(nDraws, shape1=alpha, shape2=beta)
}

nDraws <- 10000
alpha <- alpha_prior + s
beta <- beta_prior + f

set.seed(123)
samples <- drawSamples(nDraws, alpha, beta)

true_mean <- alpha / (alpha + beta)
estimated_mean <- mean(samples)

true_sd <- sqrt((alpha * beta) / ((alpha + beta)^2 * (alpha + beta + 1)))
estimated_sd <- sd(samples)

set.seed(123)
```


```

betamean <- sapply(1:nDraws, FUN = function(X) {
  mean(drawSamples(nDraws = X, alpha = alpha , beta = beta))
})

set.seed(123)
betasd <- sapply(1:nDraws, FUN = function(X) {
  sd(drawSamples(nDraws = X, alpha = alpha , beta = beta))
})

plot(x = 1:nDraws, y = betamean, type = "l",
     xlab="# of Samples", ylab="Sample Mean")
abline(h = true_mean, col = "red", lwd = 2)

```




The plot above shows how the posterior mean changes with the number of drawn samples. The red line is the expected mean and we can see that as the number of samples increase it converges to the true mean.

```

plot(x = 1:nDraws, y = betasd, type = "l",
     xlab="# of Samples", ylab="Sample Standard Deviation")
abline(h = true_sd, col = "red", lwd = 2)

```



Similar, the plot above shows how the standard deviation from the posterior changes with the number of drawn samples. The red line shows the expected standard deviation and we can see that the posterior standard deviation converges to it.

b

```
est_04 <- sum(samples < 0.4) / length(samples)
true_04 <- pbeta(0.4, shape1 = alpha, shape2 = beta)
```

The posterior probability for theta being below 0.4 is calculated as the share of values under 0.4 in our drawn samples. We drew 10000 samples and got $\Pr(\theta < 0.4|y) \approx 0.0035$ while the true value is $\Pr(\theta < 0.4|y) = 0.0039727$ which is fairly close.


c

We can also use the samples to get the distribution of a function of those samples. In this case we take $\text{logit}(\theta) = \log(\theta/(1 - \theta))$ and get the histogram below.

```
log_odds <- function(theta){
  log(theta / (1 - theta))
}

sample_odds <- log_odds(samples)
```

```
hist(sample_odds, breaks=50, freq = FALSE, main="Histogram of Log Odds", xlab="Log Odds")  
lines(density(sample_odds), col = "blue")
```



Question 2

a

```
library(geoR)
sek <- c(14, 25, 45, 25, 30, 33, 19, 50, 34 ,67)

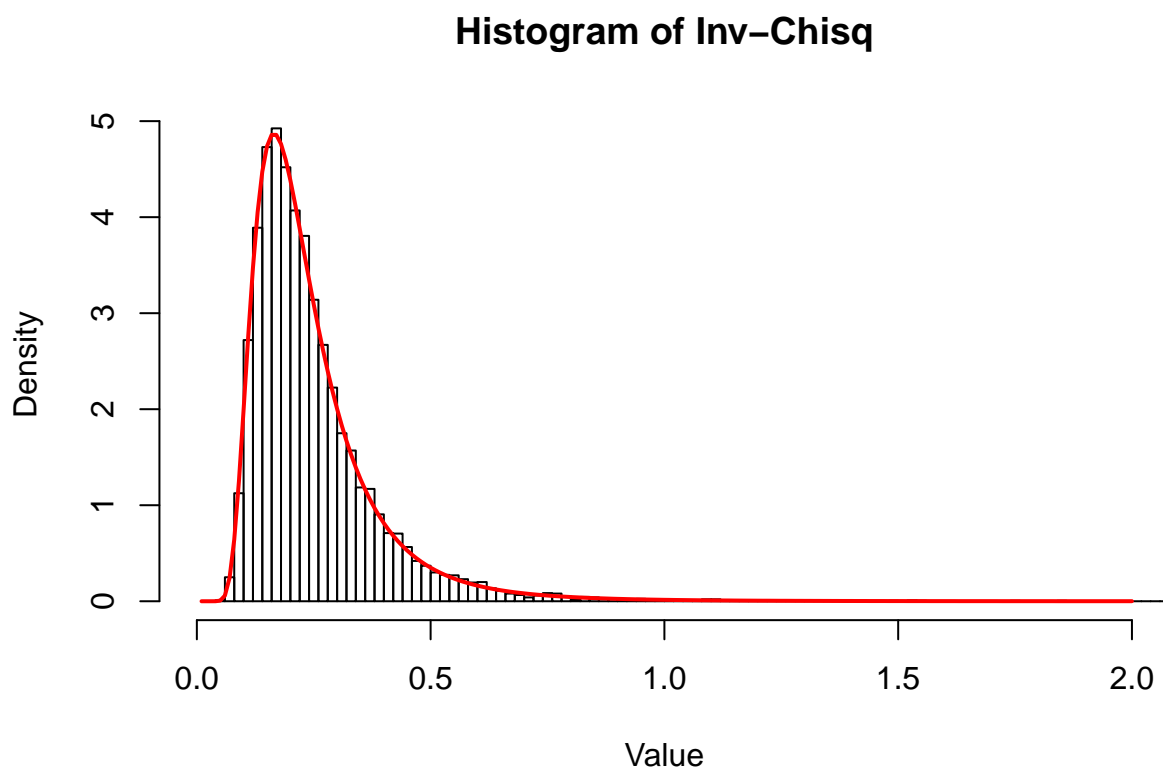
n <- length(sek)
mu <- 3.5
tausq <- sum((log(sek)- mu)^2) / n

set.seed(123)
chi_samples <- rinvchisq(nDraws, df = n, scale = tausq)

invchisq_pdf <- function(x, df, scale) {
  factor1 <- (scale * df / 2)^(df / 2) / gamma(df / 2)
  factor2 <- exp(-(df * scale) / (2 * x)) / x^(1 + df / 2)
  factor1 * factor2
}

xs <- seq(0.01, 2, 0.01)
ys <- invchisq_pdf(xs, 10, tausq)


hist(chi_samples, breaks = 100, freq=FALSE, xlim=c(0, 2),
     main="Histogram of Inv-Chisq", xlab="Value")
lines(xs, ys, col="red", lwd=2)
```



The red line represents the theoretical distribution and it looks like a really nice fit to the posterior in the histogram.

b

```
gini <- function(sigma){  
  2 * pnorm(sigma / sqrt(2), mean = 0 , sd = 1) - 1  
}  
  
gini_samples <- gini(sqrt(chi_samples))  
  
hist(gini_samples, breaks = 50, freq = FALSE,  
     ylim = c(0,8), xlim=c(0, 1), main="Histogram of Gini Coefficient", xlab="Value")
```



The histogram above shows the posterior distribution of the Gini coefficient.

c

```
gini_cred <- quantile(gini_samples, probs = c(0.025, 1 - 0.025))

dense_gini <- density(gini_samples)
dense_gini_cdf <- dense_gini$y/sum(dense_gini$y)
index_gini <- order(dense_gini_cdf)
dense_gini_cdf <- dense_gini_cdf[index_gini]


while(sum(dense_gini_cdf) > 0.95){
  dense_gini_cdf <- dense_gini_cdf[-1]
  index_gini <- index_gini[-1]
}

lower <- min(dense_gini$x[index_gini])
upper <- max(dense_gini$x[index_gini])

gini_HPD <- c(lower, upper)

hist(gini_samples, breaks = 50, freq = FALSE, ylim = c(0, 8),
     main="Histogram of Gini Coefficient", xlab="Value")
lines(density(gini_samples), col = "red", lwd = 2)
abline(v = gini_HPD, col = "blue", lwd = 2)
```

```
abline(v = gini_cred, col = "orange", lwd = 2, lty = 2)
```



The credible-interval is represented by the dashed orange line and the parameter G lies within this interval with 0.95 probability and the probability of G being outside the interval is equally likely on both sides.

The Highest Posterior Density (HPD) interval contains 95% of the most probable G values.

The main difference between the two intervals is that the HPD is more shifted to the left since the distributions is positively skewed.

What we can say about the data is that it is more probable that the income-distribution is equal than unequal.

Question 3

For this exercise we have been given the likelihood

$$p(y|\mu, \kappa) = \frac{\exp[\kappa \cos(y - \mu)]}{2\pi I_0(\kappa)}, -\pi \leq y \leq \pi$$

where $I_0(\kappa)$ is the modified Bessel function of the first kind of order zero. We have also been given that $\kappa \sim \text{Exponential}(\lambda = 1)$. We want to approximate the marginal posterior distribution $p(\kappa|\mu, y)$ and we assume that $\mu = 2.39$. In order to this we generate samples of κ that we assume comes from the Exponential distribution, in our case 1000 values equally spaces between 0.01 to 10. We can then compute

$$p(\kappa|\mu, y) \propto p(y|\mu, \kappa)p(\kappa)$$

and get an approximation of the marginal posterior distribution. This method usually goes by the name of *grid approximation* and is appropriate for low number of parameters.

a

```
vonMises <- function(y,K,mu){
  exp(K * cos(y - mu)) / (2 * pi * besseli(K, nu = 0))
}


y <- c(-2.44, 2.14, 2.54, 1.83, 2.02, 2.33, -2.79, 2.23, 2.07, 2.02)

sample_exp <- seq(0.01, 10, 0.01)
sample_mises <- t(sapply(y, function(X) vonMises(y = X, K = sample_exp, mu = 2.39)))

prior <- dexp(sample_exp, rate = 1)
likelihood <- apply(sample_mises, MARGIN = 2, FUN = prod)

unnorm_posterior <- likelihood * prior
posterior <- unnorm_posterior / sum(unnorm_posterior)

plot(sample_exp, posterior, type="l", ylab="Density", xlab="k")
```



We can see from the distribution above that the most density is located between 0.5 and 4.

b

The posterior mode of κ is approximately 2.12 given by our approximation of the distribution.