

Bayesian Learning

Lab 3

Emil K Svensson and Rasmus Holm

2017-05-11

Question 1

```
options(digits=2)

rainfall <- read.table("../data/rainfall.dat", header=F)
```

Prior

$$\begin{aligned}\mu &\sim N(\mu_0, \tau_0^2) \\ \sigma^2 &\sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)\end{aligned}$$

Likelihood

$$\mathbf{y}|\mu, \sigma^2 \sim N(\mu, \sigma^2)$$

Posterior

$$\begin{aligned}\mu|\sigma^2, \mathbf{x} &\sim N(\mu_n, \tau_n^2) \\ \sigma^2|\mu, \mathbf{x} &\sim \text{Inv-}\chi^2\left(\nu_n, \frac{\nu_0\sigma_0^2 + \sum_{i=1}^n (x_i - \mu)^2}{n + \nu_0}\right)\end{aligned}$$

where

$$\begin{aligned}\mu_n &= \frac{\frac{1}{\tau_0^2}}{\frac{1}{\tau_0^2} + \frac{n}{\sigma^2}} \mu_0 + \frac{\frac{n}{\sigma^2}}{\frac{1}{\tau_0^2} + \frac{n}{\sigma^2}} \bar{x} \\ \frac{1}{\tau_n^2} &= \frac{1}{\tau_0^2} + \frac{n}{\sigma^2} \\ \nu_n &= \nu_0 + n\end{aligned}$$

a)

```
library(geoR)

mun <- function(x, sigmasq, hyperparams){
  n <- length(x)
```

```

    denom <- ((1 / hyperparams$tausq0) + (n / sigmasq))

    pt1 <- ((1 / hyperparams$tausq0) / denom) * hyperparams$mu0
    pt2 <- ((n / sigmasq) / denom) * mean(x)

    pt1 + pt2
}

taun <- function(x, sigmasq, hyperparams){
  hyperparams$tausq0 + (sigmasq / length(x))
  1 / hyperparams$tausq0
}

nun <- function(x, hyperparams){
  hyperparams$nu0 + length(x)
}

sigmasqn <- function(x, mu, hyperparams){
  (hyperparams$nu0 * hyperparams$sigmasq0 + sum((x - mu)^2 )) / (length(x) + hyperparams$nu0)
}

musampler <- function(x,sigmasq, hyperparams){
  mu <- mun(x, sigmasq, hyperparams)
  sigma <- sqrt(taun(x,sigmasq, hyperparams))
  rnorm(1, mu, sigma)
}

sigmasampler <- function(x, mu, hyperparams){
  scale <- sigmasqn(x, mu, hyperparams)
  df <- nun(x, hyperparams)
  rinvcchisq(1, df, scale)
}

gibbs <- function(x, iter, init, hyperparams){
  samples <- matrix(NA, ncol = 2, nrow = iter + 1)
  samples[1,] <- init

  for (i in 2:(iter+1)){
    mu <- musampler(x, samples[i-1, 2], hyperparams)
    sigma <- sigmasampler(x, mu, hyperparams)
    samples[i,] <- c(mu, sigma)
  }

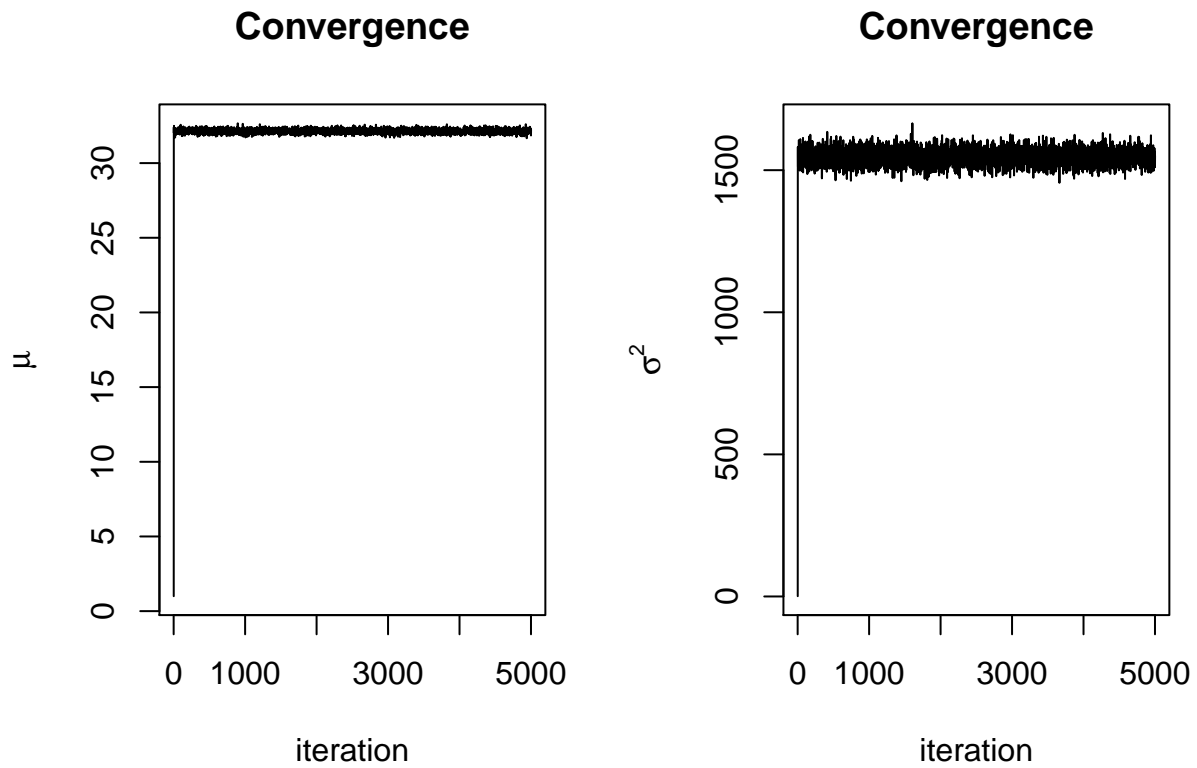
  colnames(samples) <- c("mu", "sigmasq")
  samples
}

set.seed(123456)

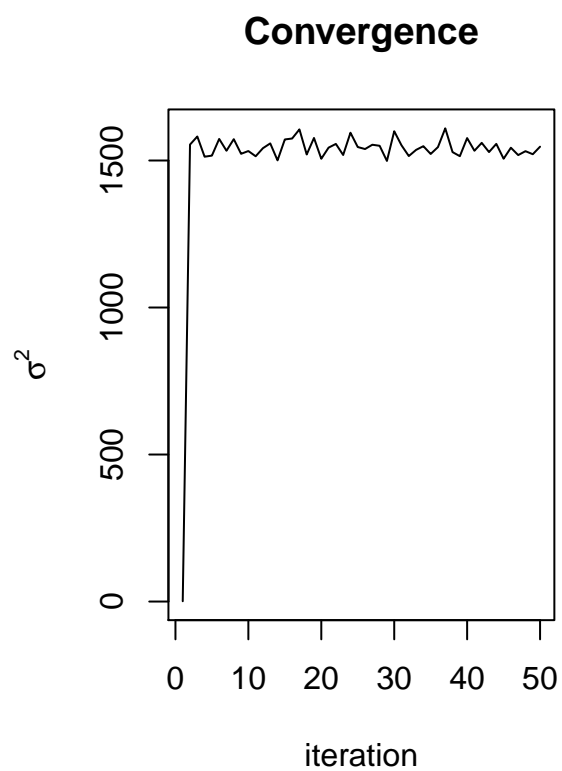
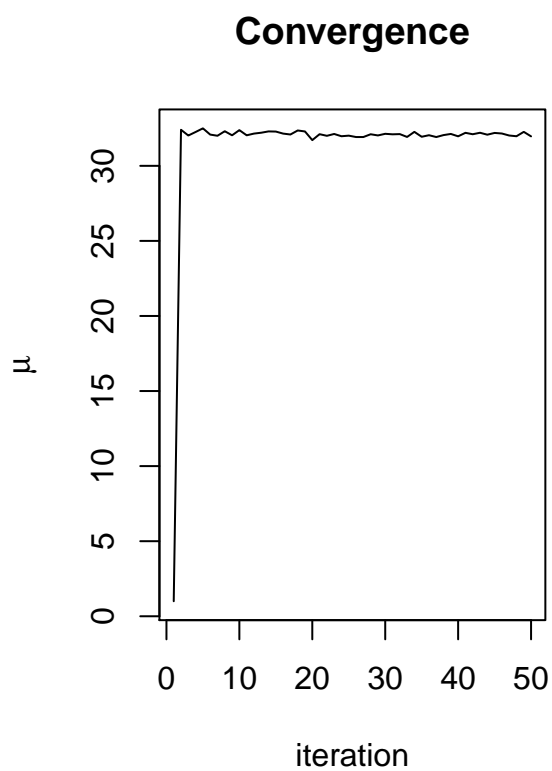
hyperparams <- list(mu0=0, tausq0=50, nu0=1, sigmasq0=50)
iter <- 5000

```

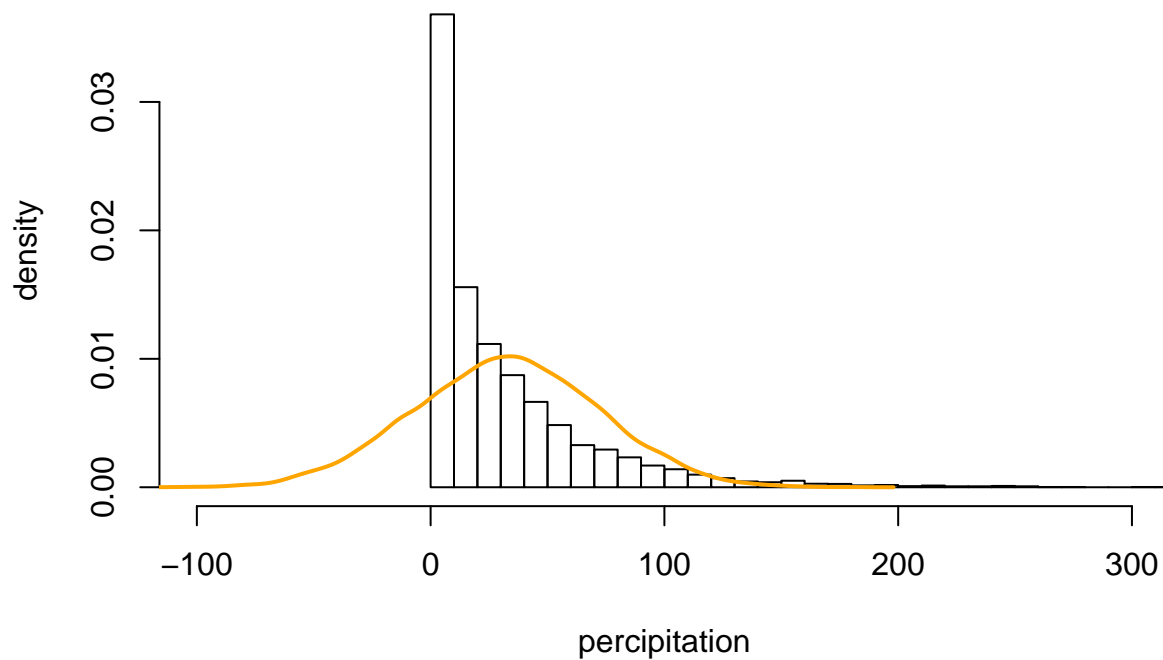
```
gibbs_params <- gibbs(x=rainfall$V1, iter = iter, init = c(1,1), hyperparams)
```



The μ converges to 32.13 which is similar to the data mean, 32.28, and σ^2 converges to 1546.15 which is similar to data variance, 1547.1.



The convergence happens after just a few iterations.



The data is obviously not normally distributed so the model is not particularly good which is seen by the blue curve.

b)

```
nn <- function(I){
  n1 <- sum(I == 1)
  n2 <- sum(I == 2)
  list(n1 = n1, n2 = n2)
}

dd <- function(x, I) {
  d1 <- x[I == 1]
  d2 <- x[I == 2]
  list(d1=d1, d2=d2)
}

pisampler <- function(x, I, hyperparams){
  n <- nn(I)
  rbeta(1, shape1 = n$n1 + hyperparams$a1, shape2 = n$n2 + hyperparams$a2)
}

sigmasq2sampler <- function(x, I, hyperparams){
  n <- nn(I)
  d <- dd(x, I)
```

```

vnsn1 <- hyperparams$nu0 * hyperparams$sigmasq0 +
  (n$n1 - 1) * var(d$d1) +
  (((1 / hyperparams$tausq0) * n$n1) / ((1 / hyperparams$tausq0) + n$n1)) *
  ((mean(d$d1) - hyperparams$mu0)^2)

vnsn2 <- hyperparams$nu0 * hyperparams$sigmasq0 +
  (n$n2 - 1) * var(d$d2) +
  (((1 / hyperparams$tausq0) * n$n2) / ((1 / hyperparams$tausq0) + n$n2)) *
  ((mean(d$d2) - hyperparams$mu0)^2)

vn1 <- hyperparams$nu0 + n$n1
vn2 <- hyperparams$nu0 + n$n2

sn1 <- vnsn1 / vn1
sn2 <- vnsn2 / vn2

sigmasq1 <- rinvchisq(1, vn1, sn1)
sigmasq2 <- rinvchisq(1, vn2, sn2)

c(sigmasq1, sigmasq2)
}

mu2sampler <- function(x, I, sigmasq, hyperparams){
  d <- dd(x, I)

  sigma1 <- sqrt(taun(d$d1, sigmasq[1], hyperparams))
  mu1 <- mun(d$d1, sigmasq[1], hyperparams)

  sigma2 <- sqrt(taun(d$d2, sigmasq[2], hyperparams))
  mu2 <- mun(d$d2, sigmasq[2], hyperparams)

  mu1 <- rnorm(1, mu1, sigma1)
  mu2 <- rnorm(1, mu2, sigma2)

  c(mu1, mu2)
}

Isampler <- function(x, pi, mu, sigmasq){
  nom <- (1 - pi) * dnorm(x, mu[2], sqrt(sigmasq[2]))
  denom <- pi * dnorm(x, mu[1], sqrt(sigmasq[1])) + nom
  theta <- nom / denom
  rbinom(length(x), prob = theta, size = 1) + 1 # to get them into 1 and 2
}

ysampler <- function(pi, mu, sigmasq){
  pi * rnorm(1, mu[1], sqrt(sigmasq[1])) +
  (1 - pi) * rnorm(1, mu[2], sqrt(sigmasq[2]))
}

mixedgibbs <- function(x, iter, init, hyperparams){
  params_samples <- matrix(NA, ncol = 5, nrow = iter)
  samples <- rep(NA, iter)
  I <- init

```

```

for (i in 1:iter){
  pi <- pisampler(x, I, hyperparams)
  sigmasq <- sigmasq2sampler(x, I, hyperparams)
  mu <- mu2sampler(x, I, sigmasq, hyperparams)
  I <- Isampler(x, pi, mu, sigmasq)

  samples[i] <- ysampler(pi, mu, sigmasq)
  params_samples[i,] <- c(mu, sigmasq, pi)
}

colnames(params_samples) <- c("mu1", "mu2", "sigmasq1", "sigmasq2", "pi")
list(samples=samples, params=params_samples)
}

```

```
set.seed(123456)
```

```


hyperparams <- list(a1=1, a2=1, mu0=1, tausq0=50, nu0=1, sigmasq0=50)
init <- sample(c(1, 2), size=length(rainfall$V1), replace=TRUE)
iter <- 10000



```


```

mixed <- mixedgibbs(rainfall$V1, iter, init, hyperparams)
mixed_params <- mixed$params
mixed_samples <- mixed$samples

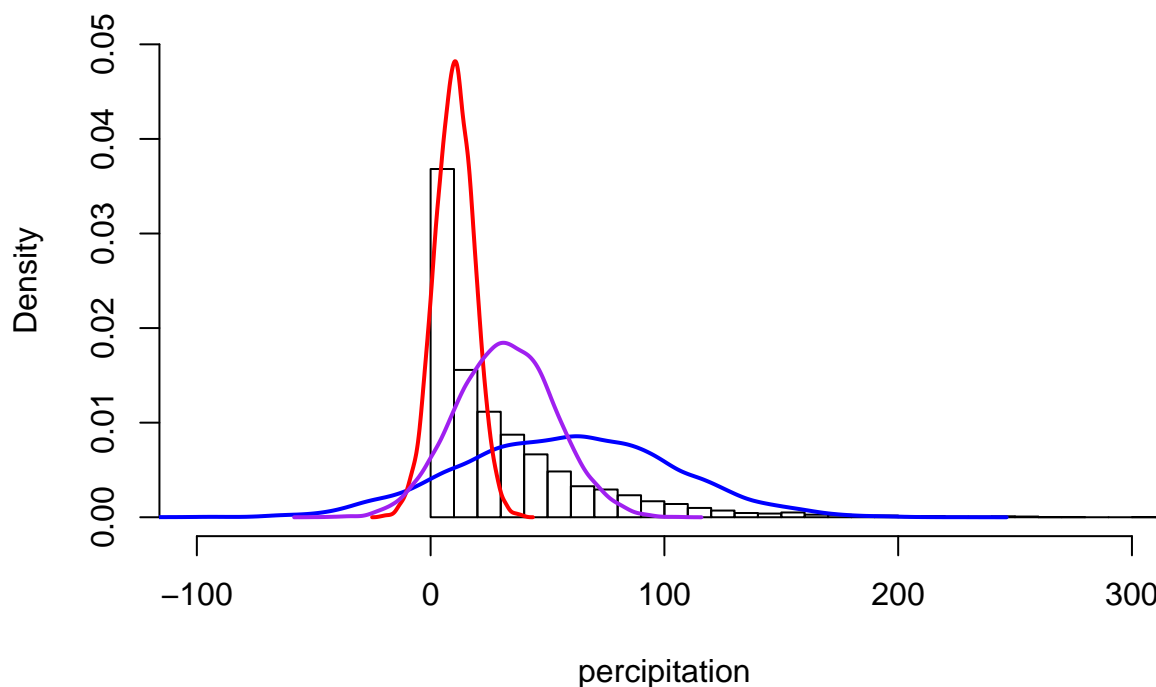
```







Densities of Normals



We can see that the two Gaussian distributions in the mixture model represent low values with high density, the red curve, and those with higher values with a wider distribution, the blue curve. Then the model, the purple curve, is an average of those aforementioned distributions. However, as mention before the data is not normally distributed so the final model is far from perfect.

c)

```
n <- 10000
burnins <- 1000 + 1

gibbs_params <- gibbs_params[-(1:burnins),]
mixed_params <- mixed_params[-(1:burnins),]

hist(rainfall$V1, freq=FALSE, breaks=50, xlim=c(-100, 200), main="", xlab="percipitation")

## Model A
mu <- mean(gibbs_params[, "mu"])
sigmasq <- mean(gibbs_params[, "sigmasq"])

samples <- rnorm(n, mu, sqrt(sigmasq))

lines(density(samples), lwd=2, col="orange")

## Model B
mu1 <- mean(mixed_params[, "mu1"])
```

```

sigmasq1 <- mean(mixed_params[, "sigmasq1"])

mu2 <- mean(mixed_params[, "mu2"])
sigmasq2 <- mean(mixed_params[, "sigmasq2"])

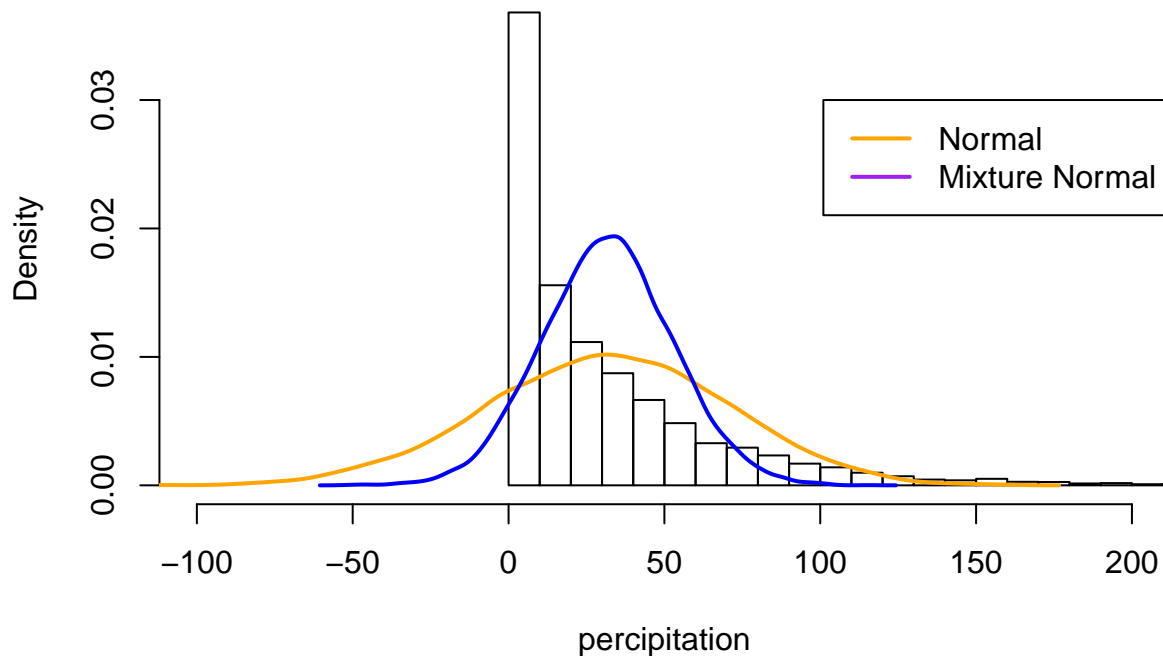
pi <- mean(mixed_params[, "pi"])

samples <- pi * rnorm(n, mu1, sqrt(sigmasq1)) + (1 - pi) * rnorm(n, mu2, sqrt(sigmasq2))

lines(density(samples), lwd=2, col="blue")

legend(101, 0.03, lty=c(1, 1), lwd=c(2, 2), col=c("orange", "purple"),
      legend=c("Normal", "Mixture Normal"))

```



Comparing the normal model and the mixture of normals we conclude that the latter is better in this case but it would probably have been better to choose either a truncated Normal or exponential distribution as the model since we know that precipitation cannot be negative.

Question 2

```
women <- read.table("../data/WomenWork.dat", header = T)
women <- as.matrix(women)
```

a)

```
library(msm)
library(mvtnorm)

betahat <- function(XX, X, y){
  solve(XX) %*% t(X) %*% y
}

mun <- function(XX, betahat, mu0, omega0){
  solve( XX + omega0 ) %*% (XX %*% betahat - omega0 %*% mu0)
}

probitgibbs <- function(X, Y, mu0, omega0, beta0, iter){
  samples <- matrix(ncol = ncol(X), nrow = iter)
  y0 <- Y == 0
  y0n <- sum(y0)
  y1n <- sum(!y0)
  XX <- t(X) %*% X
  omegan <- XX + omega0
  u <- c(rep(1, times = length(Y)))
  beta <- beta0

  for (i in 1:iter){

    xB <- X %*% beta

    u[y0] <- rtnorm(n = y0n, mean = xB[y0], sd = 1, lower = -Inf, upper = 0)
    u[!y0] <- rtnorm(n = y1n, mean = xB[!y0], sd = 1, lower = 0, upper = Inf)

    baehat <- betahat(XX, X, u)
    myn <- mun(XX, baehat, mu0, omega0)

    beta <- as.vector(rmvnorm(1, mean = myn, sigma = 1 * solve(omegan) ))
    samples[i,] <- beta
  }
  samples
}
```

b)

```
betas <- probitgibbs(X = women[, -1], Y = women[, 1],
                    mu0 = rep(0, times = ncol(women[, -1])),
                    omega0 = 100 * diag(ncol(women[, -1])),
```

```

        beta = rep(0, times = ncol(women[,-1])), iter = 10000)
probitbetas <- colMeans(betas)

```

c)

```

logprior <- function(beta, mean, sigma){
  dmnorm(beta, mean = mean, sigma = sigma, log = TRUE)
}

loglikelihood <- function(beta, X, Y){
  linear_prediction <- t(X) %*% beta
  P <- pnorm(linear_prediction, mean = 0, sd = 1)
  probabilities <- Y * log(P) + (1 - Y)*log(1 - P)
  loglike <- sum(probabilities)

  ## if (abs(loglike) == Inf)
  ##   loglike = -20000

  loglike
}

logposterior <- function(beta, X, Y, mean, sigma){
  loglikelihood(beta, X, Y) + logprior(beta, mean, sigma)
}

tau <- 10
mu <- rep(0,8)
sigma <- tau^2 * diag(8)

womenX <- as.matrix(women[,2:ncol(women)])
womenY <- as.matrix(women[,1])

quadapprox <- optim(par = matrix(rep(0, 8), ncol = 1),
  fn = logposterior, method = "BFGS", hessian = TRUE,
  X = t(womenX), Y = womenY,
  mean = mu, sigma = sigma,
  control=list(fnscale=-1))

options(digits=3)

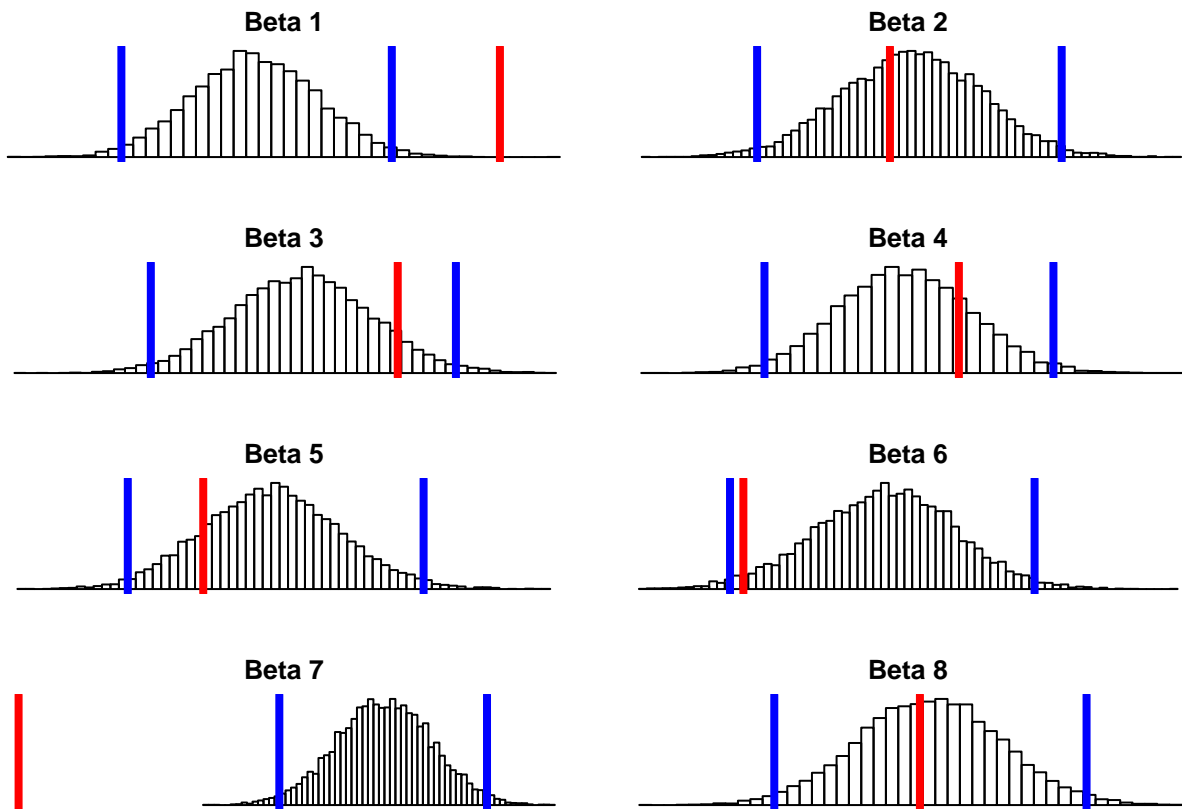
probitbetas

## [1] -0.00572 -0.00944  0.06510  0.08245 -0.00918 -0.03144 -0.15049 -0.00455

quadbetas <- as.vector(quadapprox$par)
quadbetas

## [1]  0.3845 -0.0121  0.1084  0.1022 -0.0902 -0.0496 -0.8157 -0.0124

```



The blue lines show the 97% credibility intervals and the red lines represent the corresponding parameter value found by normal approximation.